

# Weather in Australia - Team 7

This cell just loads all used modules for running the notebook. Please install any package if you don't have it installed in your environment so far.

```
In [1]: #disable some annoying warnings
import warnings
warnings.filterwarnings('ignore', category=FutureWarning)
#-----#
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import pyplot
#plots the figures in place instead of a new window
%matplotlib inline

import statistics

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn import decomposition
from numpy import unique
from numpy import where
from sklearn.datasets import make_classification
from sklearn.cluster import KMeans
from matplotlib import pyplot
from sklearn.cluster import AffinityPropagation
from sklearn.cluster import AgglomerativeClustering
from IPython.display import display, clear_output

from sklearn.ensemble import GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn import tree
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, confusion_matrix, recall_score, precision_score, explained_variance_score, mean_squared_error, r2_score, mean_absolute_error
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
from xgboost import XGBClassifier, XGBRegressor

from abess import LinearRegression
import statsmodels.api as sm
```

## Dataset Overview

We chose the rain in Australia dataset from Kaggle because we thought that it could be interesting to analyze a dataset with around 145000 rows. It is also interesting that data from about 10 years of daily observations from different locations throughout Australia has been collected.

Besides several numerical attributes, also several categorical attributes are provided. The attributes of the used dataset are explained below.

1. Date: The observation's date
2. Location: The location of the observation
3. MinTemp: The minimum temperature on that day (°C)
4. MaxTemp: The maximum temperature on that day (°C)
5. Rainfall: The rainfall amount measured in mm
6. Evaporation: The evaporation also measured in mm
7. Sunshine: The number of sunshine hours
8. WindGustDir: The strongest wind gust's direction
9. WindGustSpeed: The strongest wind gust's speed in km/h
10. WindDir9am: The wind's direction at 9 AM
11. WindDir3pm: The wind's direction at 3 PM
12. WindSpeed9am: The wind's speed (km/h) at 9 AM
13. WindSpeed3pm: The wind's speed (km/h) at 3 PM
14. Humidity9am: The humidity percentage at 9 AM
15. Humidity3pm: The humidity percentage at 3 PM
16. Pressure9am: The atmospheric pressure (hpa) at 9 AM
17. Pressure3pm: The atmospheric pressure (hpa) at 3 PM
18. Cloud9am: Fraction of obscured sky by clouds (in "oktas") at 9 AM
19. Cloud3pm: Same as above but at 3 PM
20. Temp9am: Temperature in °C at 9 AM
21. Temp3pm: Temperature in °C at 3 PM
22. RainToday: True, if it has been raining on that day, otherwise False
23. RainTomorrow: True, if it has been raining on the next day, otherwise False; target variable

In [2]: 

```
# use the weather dataset of heterogenous data and plot first 5 lines
weather = pd.read_csv('data/weatherAUS.csv')
weather.head()
```

Out[2]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	44.0	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	44.0	I
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	46.0	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	24.0	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	41.0	

5 rows × 23 columns

In [3]: 

```
# overview of the created datatypes
weather.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Date                145460 non-null object
```

```

1 Location      145460 non-null object
2 MinTemp      143975 non-null float64
3 MaxTemp      144199 non-null float64
4 Rainfall     142199 non-null float64
5 Evaporation  82670 non-null float64
6 Sunshine     75625 non-null float64
7 WindGustDir  135134 non-null object
8 WindGustSpeed 135197 non-null float64
9 WindDir9am   134894 non-null object
10 WindDir3pm  141232 non-null object
11 WindSpeed9am 143693 non-null float64
12 WindSpeed3pm 142398 non-null float64
13 Humidity9am  142806 non-null float64
14 Humidity3pm  140953 non-null float64
15 Pressure9am  130395 non-null float64
16 Pressure3pm  130432 non-null float64
17 Cloud9am     89572 non-null float64
18 Cloud3pm     86102 non-null float64
19 Temp9am      143693 non-null float64
20 Temp3pm      141851 non-null float64
21 RainToday    142199 non-null object
22 RainTomorrow 142193 non-null object
dtypes: float64(16), object(7)
memory usage: 25.5+ MB

```

## Data Preparation - Adjust Date Values

In this step, the data gets adjusted, in order to fit for our analysis. This adjustments go especially for the Date in the first place. Here the whole Date value gets split up into a new year month and day column, in order to better aggregate over the set.

```

In [4]: # Convert Date to a date type and create new columns
weather['Date_converted'] = pd.to_datetime(weather['Date'], format='%Y-%m-%d')
weather['Year'] = weather['Date_converted'].dt.year
weather['Month'] = weather['Date_converted'].dt.month
weather['Day'] = weather['Date_converted'].dt.day

```

## Overview of missing values


In order to do a proper data cleaning and having a feeling, how many values are even missing, we analysed the amount of missing data per column. It can be seen that for some columns nearly half of the values (40 - 48%) are missing (shown in the table as well as the plot above).

```

In [5]: # Calculate percentage of null values per attribute
missing_in_percentage = weather.isnull().sum() * 100 / len(weather)
missing = pd.DataFrame({'col': weather.columns, 'missing_percent': missing_in_percentage})
missing.sort_values('missing_percent', inplace=True, ascending=False)

ax = sns.barplot(x="col", y="missing_percent", data=missing.head(8))
ax.set_ylim((0, 100))
ax.set_xticklabels(ax.get_xticklabels(), rotation=300)
ax.set_title('Percentage of missing values per DF column')
ax.set_xlabel('DF columns')
_ = ax.set_ylabel('Percentage of missing values')

```



## Base for missing values

### Missing values in different seasons

Now we further investigate this issue by looking at the columns sunshine, evaporation, cloud3pm and cloud9am by grouping the percentage of missing values first by season, to look whether we can see a seasonal affect. We also group the percentage of missing values by location to see if we can spot a locational affect. But as you can also see in the table below, there is no real trend, if the values tend to be not recorded in a specific season.

```
In [6]: # Mapping the dates to seasons and calculate for each season and attribute the percentage
seasons = {
    1: 'Winter',
    2: 'Spring',
    3: 'Summer',
    4: 'Autumn'
}
df_values_season = weather[['Year', 'Month', 'Sunshine', 'Evaporation', 'Cloud3pm', 'Cloud9am']]
df_values_season['Season'] = (df_values_season['Month'] % 12 + 3) // 3
df_values_season['Season_name'] = df_values_season['Season'].map(seasons)

df_season_count_null = df_values_season[['Sunshine', 'Evaporation', 'Cloud3pm', 'Cloud9am']]
df_season_count_all = df_values_season[['Sunshine', 'Evaporation', 'Cloud3pm', 'Cloud9am']]

df_missing_values_percent = (df_season_count_null / df_season_count_all) * 100
df_missing_values_percent['Season'] = df_missing_values_percent.index.tolist()
df_missing_values_percent.style.hide_index()
```

```
Out[6]:
```

Sunshine	Evaporation	Cloud3pm	Cloud9am	Season
47.395082	42.394657	40.996689	38.537510	Autumn
48.680222	43.780054	41.022894	38.830232	Spring
48.109535	42.861420	38.639519	36.793968	Summer
47.793406	43.593759	42.648482	39.562098	Winter

# Missing values in different locations

As it can be seen, for 22 of the 49 locations no values are tracked which explains the large amount of missing data for the attributes 'Sunshine', 'Evaporation', 'Cloud3pm' and 'Cloud9am'. The reason for this is, however, unknown.

```
In [7]: df_values_location = weather[['Location', 'Sunshine', 'Evaporation', 'Cloud3pm', 'Cloud9am']]
df_values_location_count_null = weather[['Sunshine', 'Evaporation', 'Cloud3pm', 'Cloud9am']].isnull().sum()
# fillna is needed in order to get the
df_values_location_count_all = weather[['Sunshine', 'Evaporation', 'Cloud3pm', 'Cloud9am']].count()

df_missing_values_percent = (df_values_location_count_null / df_values_location_count_all) * 100
df_missing_values_percent['Location'] = df_missing_values_percent.index.tolist()
mask = (df_missing_values_percent == 100.).any(axis=1)
print(f'Untracked values based on location: {df_missing_values_percent[mask].shape[0]} of {df_values_location_count_all.shape[0]}')
```

Untracked values based on location: 22 of 49

## Remove missing values

Since we can not clearly 'clean' missing values in any case, because we don't have information about the geo coordinates and also no mapping of close location, we simply drop these values. Still - 112925 samples are present

```
In [8]: weather.drop(['Date', 'Sunshine', 'Evaporation', 'Cloud3pm', 'Cloud9am'], axis=1, inplace=True)
```

## Create artificial data for missing values in numeric attribute vectors when possible

For numeric data we set missing values for numeric attributes (given in the numerical\_columns value) to the median based on the year, month and (location) when possible

For the categorical values we used the mode, imputation is based on location and current month, if we do not have data for a location then only the month was used.

```
In [9]: numerical_columns = ["Pressure9am", "Pressure3pm", "Humidity3pm", "Humidity9am", "WindGustDir",
                             "WindSpeed3pm", "WindSpeed9am", "Temp9am", "MinTemp", "MaxTemp", "Rainfall"]

for col in numerical_columns:
    weather[col] = weather[col].fillna(weather.groupby(['Year', 'Month', 'Location'])[col].transform('median'))
    weather[col] = weather[col].fillna(weather.groupby(['Year', 'Month'])[col].transform('median'))

categorical_columns = ["WindDir9am", "WindGustDir", "WindDir3pm"]

for col in categorical_columns:
    weather[col] = weather[col].fillna(weather.groupby(['Year', 'Month', 'Location'])[col].transform('mode'))
    weather[col] = weather[col].fillna(weather.groupby(['Year', 'Month'])[col].transform('mode'))
```

```
In [10]: weather.dropna(inplace=True)
print(f'Amount of samples without missing values in any column: {weather.shape[0]}')
weather.head()
```

Amount of samples without missing values in any column: 140787

```
Out[10]:
```

	Location	MinTemp	MaxTemp	Rainfall	WindGustDir	WindGustSpeed	WindDir9am	WindDir3pm	WindSpeed
0	Albury	13.4	22.9	0.6	W	44.0	W	WNW	

1	Albury	7.4	25.1	0.0	WNW	44.0	NNW	WSW
2	Albury	12.9	25.7	0.0	WSW	46.0	W	WSW
3	Albury	9.2	28.0	0.0	NE	24.0	SE	E
4	Albury	17.5	32.3	1.0	W	41.0	ENE	NW

5 rows × 22 columns


## Check for valid values in all remaining (numeric) columns

In the next step, # check for minimum and maximum values in numeric attributes (in our case all attributes in the frame which have the datatype of float64. Here no out of range values could be detected.

```
In [11]: # check for minimum and maximum values in numeric attributes:
for col in weather.loc[:, weather.dtypes == 'float64']:
    print(f'Attribute {col}:')
    print("Min: {:.2f}, Q1: {:.2f}, Median {:.2f}, Q3: {:.2f}, Max: {:.2f}".format(weath
sns.boxplot(x=weather[col])
plt.title(f'Boxplot of Attribute {col}')
plt.show()
```

Attribute MinTemp:


Min: -8.50, Q1: 7.60, Median 12.00, Q3: 16.80, Max: 33.90



Attribute MaxTemp:

Min: -4.80, Q1: 17.90, Median 22.60, Q3: 28.20, Max: 48.10


Boxplot of Attribute MaxTemp



Attribute Rainfall:

Min: 0.00, Q1: 0.00, Median 0.00, Q3: 0.80, Max: 371.00


Boxplot of Attribute Rainfall



Attribute WindGustSpeed:

Min: 6.00, Q1: 31.00, Median 39.00, Q3: 46.00, Max: 135.00


Boxplot of Attribute WindGustSpeed



Attribute WindSpeed9am:

Min: 0.00, Q1: 7.00, Median 13.00, Q3: 19.00, Max: 130.00


Boxplot of Attribute WindSpeed9am



Attribute WindSpeed3pm:

Min: 0.00, Q1: 13.00, Median 19.00, Q3: 24.00, Max: 87.00


Boxplot of Attribute WindSpeed3pm



Attribute Humidity9am:

Min: 0.00, Q1: 57.00, Median 70.00, Q3: 83.00, Max: 100.00

Boxplot of Attribute Humidity9am




Attribute Humidity3pm:

Min: 0.00, Q1: 37.00, Median 52.00, Q3: 65.00, Max: 100.00




Boxplot of Attribute Humidity3pm



Attribute Pressure9am:

Min: 980.50, Q1: 1013.30, Median 1017.60, Q3: 1022.10, Max: 1041.00


Boxplot of Attribute Pressure9am



Attribute Pressure3pm:


Min: 977.10, Q1: 1010.80, Median 1015.20, Q3: 1019.68, Max: 1039.60

Boxplot of Attribute Pressure3pm

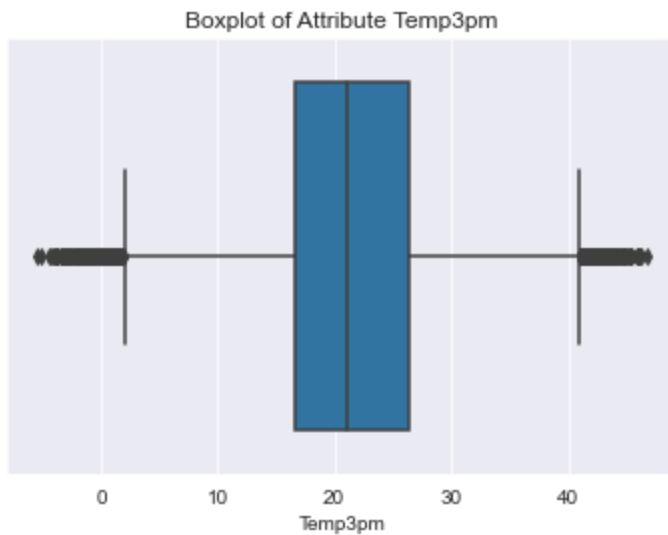


Attribute Temp9am:

Min: -7.20, Q1: 12.20, Median 16.70, Q3: 21.60, Max: 40.20



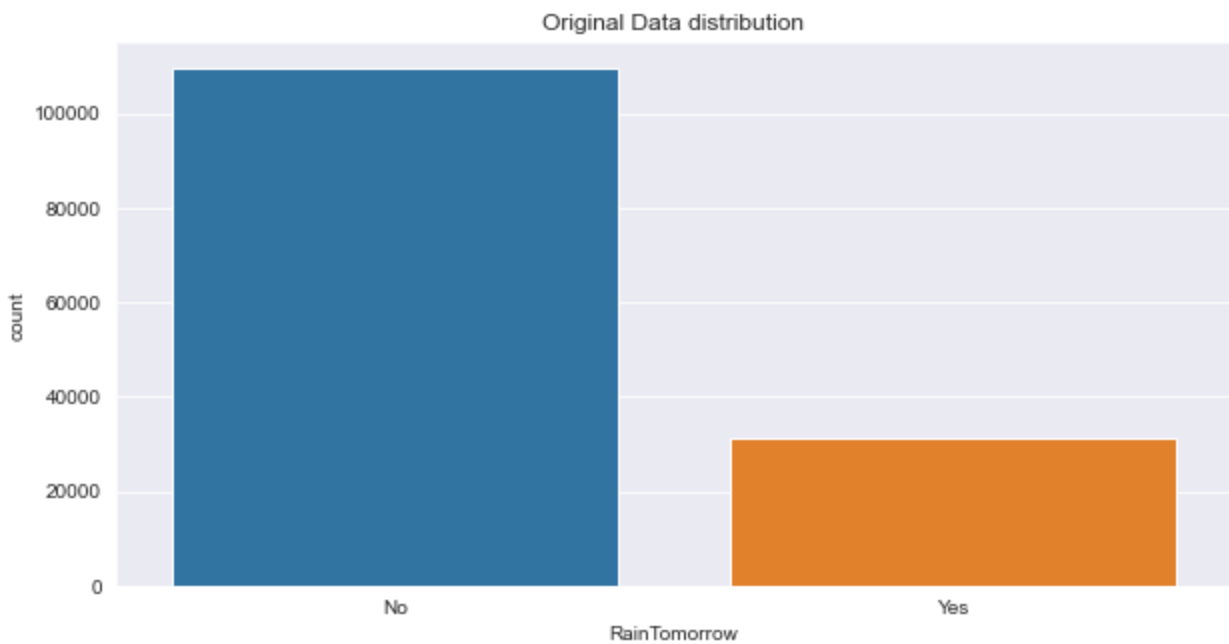
Attribute Temp3pm:  
Min: -5.40, Q1: 16.60, Median 21.10, Q3: 26.40, Max: 46.70



## Check the distribution of RainTomorrow samples

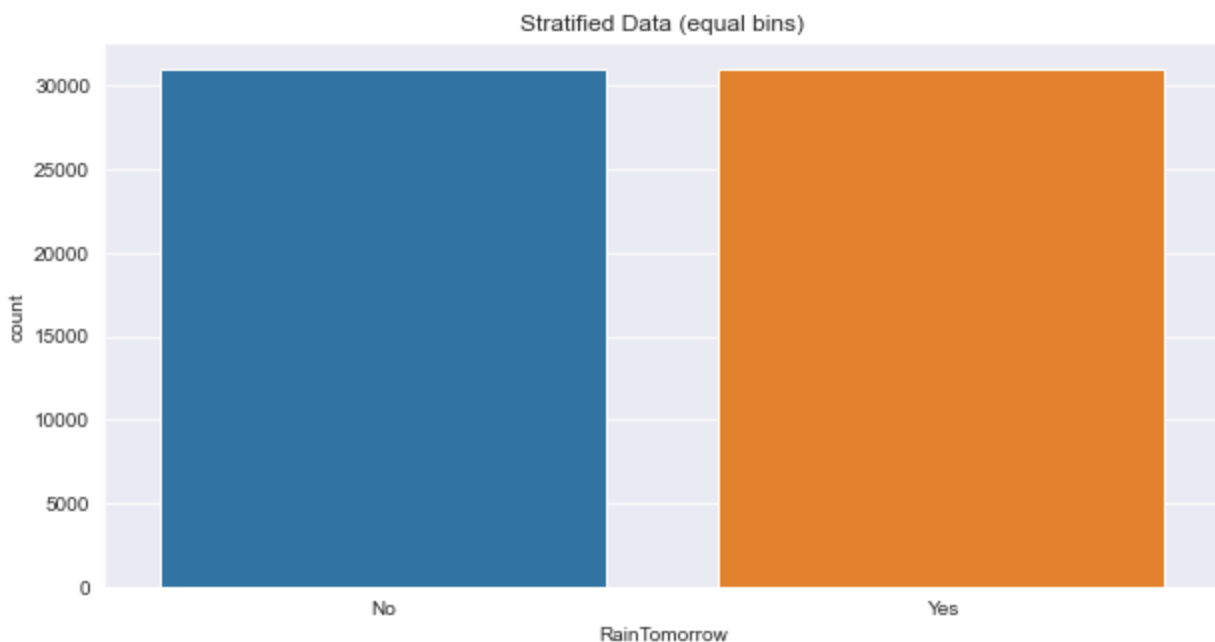
As we can clearly see in the next cell, there are a lot more samples of NOT-raining tomorrow, as samples WITH raining tomorrow

```
In [12]: plt.figure(figsize=(10,5))
sns.countplot(x="RainTomorrow", data=weather);
plt.title('Original Data distribution')
plt.show()
```



```
In [13]: # Disproportionate sampling:
# randomly select 4 samples from each stratum
stratified = weather.groupby('RainTomorrow', group_keys=False).apply(lambda x: x.sample(
```

```
In [14]: plt.figure(figsize=(10,5))
sns.countplot(x="RainTomorrow", data=stratified)
plt.title("Stratified Data (equal bins)")
plt.show()
```



## PCA to explore the underlying structure of the data

```
In [15]: stratified.drop('Date_converted',axis=1,inplace=True)
for col in stratified.loc[:, stratified.dtypes == object]:
    # creating instance of labelencoder
    labelencoder = LabelEncoder()
    # Assigning numerical values and storing in another column
    stratified[f'{col}_num'] = labelencoder.fit_transform(stratified[col])
```

```
# drop non-numeric column
stratified.drop(col,axis=1,inplace=True)
```

```
In [16]: stratified.head()
```

```
Out[16]:
```

	MinTemp	MaxTemp	Rainfall	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidit
<b>16898</b>	18.8	27.2	2.0	42.152514	0.0	19.0	87.0	59.0
<b>77422</b>	15.8	25.3	0.0	41.000000	20.0	17.0	82.0	67.0
<b>57631</b>	6.7	25.9	4.2	37.000000	9.0	15.0	100.0	37.0
<b>50011</b>	2.9	17.9	0.0	41.000000	7.0	22.0	63.0	33.0
<b>143273</b>	26.0	39.5	0.0	22.000000	0.0	0.0	57.0	43.3


5 rows × 21 columns

```
In [17]: n_components = 7

pca = decomposition.PCA(n_components=n_components)
pca_pos = pca.fit_transform(stratified)

stratified['pca1'] = pca_pos[:, 0]
stratified['pca2'] = pca_pos[:, 1]
```

```
In [18]: plt.figure(figsize=(6,6))
reducedPoints = stratified.groupby('RainTomorrow_num', group_keys=False).apply(lambda x:
sns.scatterplot(data=reducedPoints, x="pca1", y="pca2", hue="RainTomorrow_num",alpha=0.3
plt.title('PCA on the weather dataset, colored by RainTomorrow')
plt.show()
```



## Decision Tree

In this section, we try to fit a Decision Tree classifier to our data. Therefore we do a GridSearch, where we try different criterions, maximum depths of the tree and splitting methods. The trained classifier also gets

evaluated on 15% of the total data afterwards.

To keep the dataset clean, we removed all additional added attributes, we used in the previous section due to have more comfort. This does not change the actual data at all.

Note, that the data is also stratified like in the PCA above, so all classes are evenly distributed (standard would be to have a much higher amount of samples in the RainTomorrow=No compared to RainTomorrow=Yes)

After creating the training and test sets, training and evaluating using a confusion matrix and accuracy as a score, we also provided an overview of the feature importance learned by the decision tree.

```
In [19]: """
Evaluates the model and returns accuracy as well as a confusion matrix. Also the time for
@param model, sklearn model, trained model
@param x_test, np ndarray, data matrix
@param y_test, np ndarray, data vector
"""
def get_evaluation(model, x_test, y_test):
    y_pred = model.predict(x_test)
    accuracy = accuracy_score(y_test, y_pred)
    conf_mat = confusion_matrix(y_test, y_pred)
    rec_result = recall_score(y_test, y_pred, average=None, labels=[0,1])
    prec_result = precision_score(y_test, y_pred, average=None, labels=[0,1])

    print('\nAccuracy of Classifier on Test Image Data: ', accuracy)
    print()
    print('Recall (No Rain Tomorrow) of Classifier on Test Image Data: ', rec_result[0])
    print('Recall (Rain Tomorrow) of Classifier on Test Image Data: ', rec_result[1])
    print()
    print('Precision (No Rain Tomorrow) of Classifier on Test Image Data: ', prec_result[0])
    print('Precision (Rain Tomorrow) of Classifier on Test Image Data: ', prec_result[1])
    print()
    print('\nConfusion Matrix: \n', conf_mat)

    plt.matshow(conf_mat)
    plt.title('Confusion Matrix')
    plt.colorbar()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    return None
```

```
In [20]: def get_ROC(model, x_test, y_test):
    """
    Calculates AUC score and plots ROC curve
    @param model, sklearn model, trained model
    @param x_test, np ndarray, data matrix
    @param y_test, np ndarray, data vector
    """
    predictions = model.predict_proba(x_test)

    print('AUC score:')
    print(roc_auc_score(y_test, predictions[:,1]))

    fpr, tpr, _ = roc_curve(y_test, predictions[:,1])

    plt.clf()
    plt.plot(fpr, tpr)
    plt.xlabel('FPR')
    plt.ylabel('TPR')
```

```
plt.title('ROC curve with AUC: {:.3f}'.format(roc_auc_score(y_test, predictions[:,1]))
plt.show()
```

```
In [21]: param_grid = {
    'criterion': ['gini','entropy'],
    'max_depth': range(1,20),
    'splitter': ['random', 'best']
}

"""
Trains a decision tree using cross-validation and returns certain attributes of the rece
parameter combination.
@param x_train, np ndarray, data matrix
@param y_train, np ndarray, data vector
@param param_grid, dict, grid holding the paramaters for search
"""

def train_dec_tree(x_train,y_train,param_grid):
    tree = DecisionTreeClassifier(random_state=55)
    model = GridSearchCV(tree,param_grid=param_grid,n_jobs = -1)
    model.fit(x_train,y_train)
    return model.best_params_,model.best_estimator_
```

```
In [22]: # remove target value and additional added columns
X = stratified.drop(['RainTomorrow_num','pca1','pca2'], axis=1)
y = stratified['RainTomorrow_num']
print(f'shape of data matrix: {X.shape}')
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=2)
print(f'shape of train matrix: {x_train.shape}')
print(f'shape of test matrix: {x_test.shape}')
X.head()
```

```
shape of data matrix: (62000, 20)
shape of train matrix: (49600, 20)
shape of test matrix: (12400, 20)
```

```
Out[22]:
```

	MinTemp	MaxTemp	Rainfall	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidit
<b>16898</b>	18.8	27.2	2.0	42.152514	0.0	19.0	87.0	59.0
<b>77422</b>	15.8	25.3	0.0	41.000000	20.0	17.0	82.0	67.0
<b>57631</b>	6.7	25.9	4.2	37.000000	9.0	15.0	100.0	37.0
<b>50011</b>	2.9	17.9	0.0	41.000000	7.0	22.0	63.0	33.0
<b>143273</b>	26.0	39.5	0.0	22.000000	0.0	0.0	57.0	43.3

```
In [23]: # train decision tree with created training set and evaluate on created target set
params_dec_tree, model_dec_tree = train_dec_tree(x_train, y_train, param_grid)
_ = get_evaluation(model_dec_tree, x_test, y_test)
print("The best parameters are: {}".format(params_dec_tree))
```

```
Accuracy of Classifier on Test Image Data: 0.7587096774193548
```

```
Recall (No Rain Tomorrow) of Classifier on Test Image Data: 0.8038177735001604
```

```
Recall (Rain Tomorrow) of Classifier on Test Image Data: 0.7131041193642556
```

```
Precision (No Rain Tomorrow) of Classifier on Test Image Data: 0.7390855457227139
```

```
Precision (Rain Tomorrow) of Classifier on Test Image Data: 0.7823843416370106
```

```
Confusion Matrix:
```


```
[[5011 1223]
```

```
[1769 4397]]
```

```
The best parameters are: {'criterion': 'gini', 'max_depth': 8, 'splitter': 'best'}
```

C:\Users\fnern\AppData\Local\Temp\ipykernel\_13500\1287647186.py:27: MatplotlibDeprecatio


nWarning: Auto-removal of grids by pcolor() and pcolormesh() is deprecated since 3.5 and will be removed two minor releases later; please call grid(False) first.  
plt.colorbar()




```
In [24]: # print AUC score and ROC curve  
get_ROC(model_dec_tree, x_test, y_test)
```

AUC score:

0.8412330766242606



```
In [25]: # create overview of feature importance, of learned decision tree  
attribute_weights = pd.DataFrame({  
    'Attribute' : x_train.columns,  
    'Weight' : model_dec_tree.feature_importances_  
}).sort_values(by='Weight', ascending=False)  
plt.title('Importance of different Attributes')  
sns.barplot(data = attribute_weights, x='Weight', y='Attribute');
```



## Random Forest

```
In [26]: param_grid_forest = {
    'criterion': ['gini','entropy'],
    'max_depth': range(5,25)
}

"""
Trains a random forest using cross-validation and returns certain attributes of the rece
parameter combination.
@param x_train, np ndarray, data matrix
@param y_train, np ndarray, data vector
@param param_grid, dict, grid holding the paramaters for search
"""

def train_random_forest(x_train,y_train,param_grid):
    ensemble = RandomForestClassifier(random_state=55)
    model = GridSearchCV(ensemble,param_grid=param_grid, n_jobs = -1)
    model.fit(x_train,y_train)
    return model.best_params_,model.best_estimator_
```

```
In [27]: # train decision tree with created training set and evaluate on created target set
params_random_forest, model_random_forest = train_random_forest(x_train, y_train, param_
_ = get_evaluation(model_random_forest, x_test, y_test)
print("The best parameters are: {}".format(params_random_forest))
```

Accuracy of Classifier on Test Image Data: 0.7940322580645162

Recall (No Rain Tomorrow) of Classifier on Test Image Data: 0.8055822906641001

Recall (Rain Tomorrow) of Classifier on Test Image Data: 0.7823548491728836

Precision (No Rain Tomorrow) of Classifier on Test Image Data: 0.7891263356379635

Precision (Rain Tomorrow) of Classifier on Test Image Data: 0.7992047713717694

Confusion Matrix:


```
[[5022 1212]
```

```
[1342 4824]]
```

The best parameters are: {'criterion': 'entropy', 'max\_depth': 24}

C:\Users\fnern\AppData\Local\Temp\ipykernel\_13500\1287647186.py:27: MatplotlibDeprecatio  
nWarning: Auto-removal of grids by pcolor() and pcolormesh() is deprecated since 3.5 and  
will be removed two minor releases later; please call grid(False) first.  
plt.colorbar()






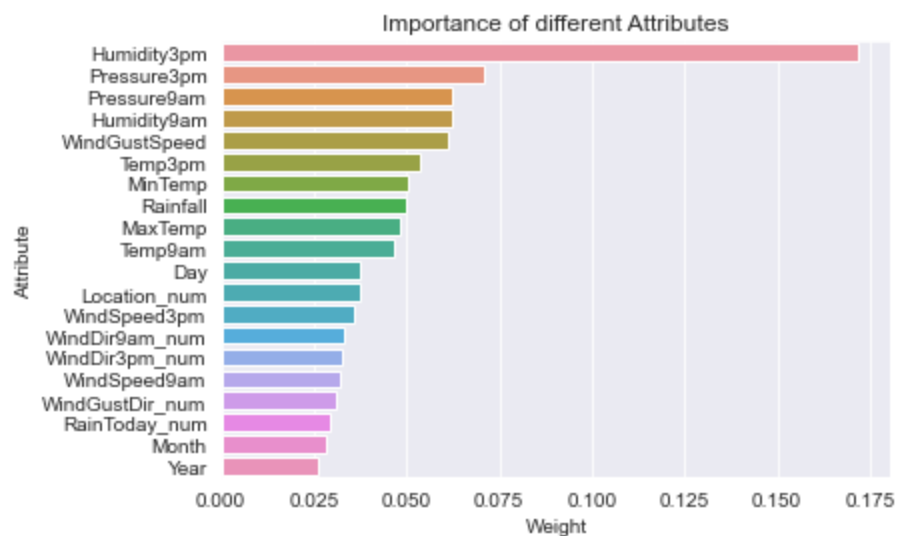
```
In [28]: # print AUC score and ROC curve
get_ROC(model_random_forest, x_test, y_test)
```

AUC score:

0.8792509733123088



```
In [29]: # create overview of feature importance, of learned decision tree
attribute_weights = pd.DataFrame({
    'Attribute' : x_train.columns,
    'Weight' : model_random_forest.feature_importances_
}).sort_values(by='Weight', ascending=False)
plt.title('Importance of different Attributes')
sns.barplot(data = attribute_weights, x='Weight', y='Attribute');
```



# Extreme Gradient Boosting

```
In [30]: xgb = XGBClassifier()  
xgb.fit(x_train, y_train)
```

C:\Users\fnern\miniforge3\envs\stat\lib\site-packages\xgboost\sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use\_label\_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num\_class - 1].

warnings.warn(label\_encoder\_deprecation\_msg, UserWarning)

[13:17:02] WARNING: D:\bld\xgboost-split\_1645118015404\work\src\learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval\_metric if you'd like to restore the old behavior.

```
Out[30]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
colsample_bynode=1, colsample_bytree=1, enable_categorical=False,  
gamma=0, gpu_id=-1, importance_type=None,  
interaction_constraints='', learning_rate=0.300000012,  
max_delta_step=0, max_depth=6, min_child_weight=1, missing=nan,  
monotone_constraints='()', n_estimators=100, n_jobs=12,  
num_parallel_tree=1, predictor='auto', random_state=0,  
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,  
tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [31]: _ = get_evaluation(xgb, x_test, y_test)
```

Accuracy of Classifier on Test Image Data: 0.7970967741935484

Recall (No Rain Tomorrow) of Classifier on Test Image Data: 0.8103946102021174

Recall (Rain Tomorrow) of Classifier on Test Image Data: 0.7836522867337009

Precision (No Rain Tomorrow) of Classifier on Test Image Data: 0.7911055433761353

Precision (Rain Tomorrow) of Classifier on Test Image Data: 0.8034585966079149


Confusion Matrix:

[[5052 1182]

[1334 4832]]

C:\Users\fnern\AppData\Local\Temp\ipykernel\_13500\1287647186.py:27: MatplotlibDeprecationWarning: Auto-removal of grids by pcolor() and pcolormesh() is deprecated since 3.5 and will be removed two minor releases later; please call grid(False) first.

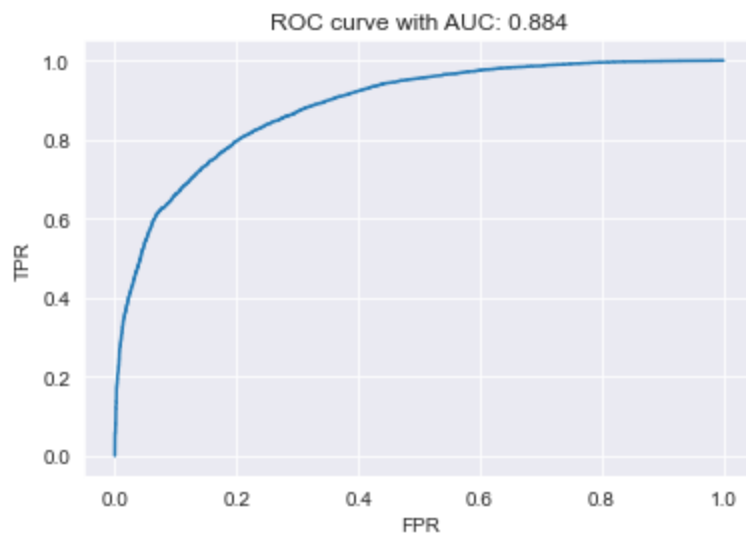
plt.colorbar()




```
In [32]: # print AUC score and ROC curve  
get_ROC(xgb, x_test, y_test)
```

AUCscore:

0.8843915285277569



```
In [33]: # create overview of feature importance, of learned decision tree
attribute_weights = pd.DataFrame({
    'Attribute' : x_train.columns,
    'Weight' : xgb.feature_importances_
}).sort_values(by='Weight', ascending=False)
plt.title('Importance of different Attributes')
sns.barplot(data = attribute_weights, x='Weight', y='Attribute');
```



## Regression

In this part, are going to develop an estimator for the rainfall. Since, rainfall is a continuous variable, this is obviously a regression task.

Since, this is going to be a multiple regression task, and therefore, not all variables might have a significant impact, we chose the best subset selection method for identifying the required variables.

## Data preparation for the regression part

```
In [34]: x_train_reg = x_train.loc[x_train['Rainfall'] > 0, x_train.columns != 'Rainfall'].copy()
x_test_reg = x_test.loc[x_test['Rainfall'] > 0, x_test.columns != 'Rainfall'].copy()

y_train_reg = x_train[x_train['Rainfall'] > 0]['Rainfall'].copy()
y_test_reg = x_test[x_test['Rainfall'] > 0]['Rainfall'].copy()
```

Now, after the data is prepared for the regression part, we can now start to fit some regression models. We decided to use the regression version of our classifiers.

Our first model is the regression tree.

## Regression tree

In [35]:

```
"""
Evaluates the regression model.
@param model, sklearn model, trained model
@param x_test, np ndarray, data matrix
@param y_test, np ndarray, data vector
"""
def get_regression_evaluation(model, x_test, y_test):
    y_pred = model.predict(x_test)

    explained_variance = explained_variance_score(y_test, y_pred)
    m_squared = mean_squared_error(y_test, y_pred)
    absolute = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    print(f"Explained variance: {explained_variance:.4f}")
    print(f"Mean squared error: {m_squared:.4f}")
    print(f"RMSE: {np.sqrt(m_squared):.4f}")
    print(f"Mean absolute error: {absolute:.4f}")
    print(f"R2 score: {r2:.4f}")

    sns.distplot(y_pred - y_test)

    return None
```

In [36]:

```
dec_tree_grid = {
    'criterion': ['squared_error', 'absolute_error'],
    'max_depth': range(1, 10),
    'splitter': ['random', 'best'],
    'max_features': ["auto", "sqrt", None],
}


"""
Trains a decision tree regressor using cross-validation and returns attributes of the re
parameter combination.
@param x_train, np ndarray, data matrix
@param y_train, np ndarray, data vector
@param param_grid, dict, grid holding the paramaters for search
@param use_pref_defined_model, bool, indicates whether the predefined model version shou
"""
def train_dec_tree_regressor(x_train, y_train, param_grid, use_pref_defined_model: bool):
    if use_pref_defined_model:
        best_params = {'criterion': 'squared_error', 'max_depth': 4, 'max_features': 'a
        tree = DecisionTreeRegressor(random_state=55, **best_params)
        tree.fit(x_train, y_train)
        return best_params, tree

    tree = DecisionTreeRegressor(random_state=55)
    model = GridSearchCV(tree, param_grid=param_grid, scoring="neg_mean_squared_error",
    model.fit(x_train, y_train)
    return model.best_params_, model.best_estimator_
```

In [37]:

```
params_dec_tree_regressor, model_dec_tree_regressor = train_dec_tree_regressor(x_train_r
_ = get_regression_evaluation(model_dec_tree_regressor, x_test_reg, y_test_reg)
#print("The best parameters are: {}".format(params_random_forest))
```

Explained variance: 0.1821  
Mean squared error: 194.2738  
RMSE: 13.9382  
Mean absolute error: 6.5251  
R2 score: 0.1821



```
In [38]: print("The best parameters are: {}".format(params_dec_tree_regressor))
```

The best parameters are: {'criterion': 'squared\_error', 'max\_depth': 4, 'max\_features': 'auto', 'splitter': 'best'}

## Random Forest Regressor

```
In [39]: rand_forest_reg_grid = {
    'n_estimators': [50, 100, 150, 200],
    'criterion': ['squared_error', 'absolute_error'],
    'max_features': ["auto", "sqrt"],
}

"""
Trains a random forest regressor using cross-validation and returns attributes of the re
parameter combination.
@param x_train, np ndarray, data matrix
@param y_train, np ndarray, data vector
@param param_grid, dict, grid holding the paramaters for search
@param use_pref_defined_model, bool, indicates whether the predefined model version shou
"""

def train_random_forest_regressor(x_train, y_train, param_grid, use_pref_defined_model:
    if use_pref_defined_model:
        best_params = {'criterion': 'squared_error', 'max_features': 'sqrt', 'n_estimato
        forest = RandomForestRegressor(random_state=55)
        forest.fit(x_train, y_train)
        return best_params, forest


    forest = RandomForestRegressor(random_state=55)
    model = GridSearchCV(forest, param_grid=param_grid, scoring="neg_mean_squared_error"
    model.fit(x_train, y_train)
    return model.best_params_, model.best_estimator_
```

```
In [40]: params_random_forest_regressor, model_random_forest_regressor = train_random_forest_regr
_ = get_regression_evaluation(model_random_forest_regressor, x_test_reg, y_test_reg)
print("The best parameters are: {}".format(params_random_forest))
```

Explained variance: 0.2415  
Mean squared error: 180.6168  
RMSE: 13.4394  
Mean absolute error: 6.4144

R2 score: 0.2396

The best parameters are: {'criterion': 'entropy', 'max\_depth': 24}



```
In [41]: _ = get_regression_evaluation(model_random_forest_regressor, x_test_reg, y_test_reg)
print("The best parameters are: {}".format(params_random_forest_regressor))
```

Explained variance: 0.2415


Mean squared error: 180.6168

RMSE: 13.4394

Mean absolute error: 6.4144

R2 score: 0.2396

The best parameters are: {'criterion': 'squared\_error', 'max\_features': 'sqrt', 'n\_estimators': 200}



## Extreme Gradient Boosting Regression

```
In [42]: xgb_grid = {
    'max_depth': [3, 6, 10],
    'learning_rate': [0.01, 0.05, 0.1],
    'n_estimators': [100, 500, 1000],
    'colsample_bytree': [0.3, 0.7]
}

"""
Trains an XGB regressor using cross-validation and returns attributes of the received mo
parameter combination.
@param x_train, np ndarray, data matrix
@param y_train, np ndarray, data vector
@param param_grid, dict, grid holding the paramaters for search
@param use_pref_defined_model, bool, indicates whether the predefined model version shou
"""
```

```

def train_xgb_regressor(x_train, y_train, param_grid, use_pref_defined_model: bool):
    if use_pref_defined_model:
        best_params = {'colsample_bytree': 0.3, 'learning_rate': 0.05, 'max_depth': 6, '
        xgb = XGBRegressor(seed = 55, **best_params)
        xgb.fit(x_train, y_train)
        return best_params, xgb
    xgb = XGBRegressor(seed = 55)
    model = GridSearchCV(xgb, param_grid=param_grid, scoring="neg_mean_squared_error", v
    model.fit(x_train, y_train)
    return model.best_params_, model.best_estimator_

```

```

In [43]: xgb_params, xgb_regressor = train_xgb_regressor(x_train_reg, y_train_reg, xgb_grid, True
_ = get_regression_evaluation(xgb_regressor, x_test_reg, y_test_reg)
print(f"The best parameters are: {xgb_params}")

```

Explained variance: 0.2652

Mean squared error: 174.5347

RMSE: 13.2112

Mean absolute error: 6.4548

R2 score: 0.2652

The best parameters are: {'colsample\_bytree': 0.3, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 500}

